


Air Quality Index



1132417385

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
plt.rcParams.update({'figure.figsize':(9,7), 'figure.dpi':120})
```

```
In [2]: data= pd.read_excel('AirQualityUCI/AirQualityUCI.xlsx', parse_dates=[['Date', 'Time']])
data.head()
```

Out[2]:

	Date_Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08.S5(O3)	T	
0	2004-03-10 18:00:00	2.6	1360.00	150	11.881723	1045.50	166.0	1056.25	113.0	1692.00	1267.50	13.60	48.
1	2004-03-10 19:00:00	2.0	1292.25	112	9.397165	954.75	103.0	1173.75	92.0	1558.75	972.25	13.30	47.
2	2004-03-10 20:00:00	2.2	1402.00	88	8.997817	939.25	131.0	1140.00	114.0	1554.50	1074.00	11.90	53.
3	2004-03-10 21:00:00	2.2	1375.50	80	9.228796	948.25	172.0	1092.00	122.0	1583.75	1203.25	11.00	60.
4	2004-03-10 22:00:00	1.6	1272.25	51	6.518224	835.50	131.0	1205.00	116.0	1490.00	1110.00	11.15	59.



In [3]: data.shape

Out[3]: (9357, 14)

In [4]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9357 entries, 0 to 9356
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date_Time             9357 non-null   datetime64[ns]
1   CO(GT)                9357 non-null   float64
2   PT08.S1(CO)           9357 non-null   float64
3   NMHC(GT)              9357 non-null   int64
4   C6H6(GT)              9357 non-null   float64
5   PT08.S2(NMHC)         9357 non-null   float64
6   NOx(GT)               9357 non-null   float64
7   PT08.S3(NOx)          9357 non-null   float64
8   NO2(GT)               9357 non-null   float64
9   PT08.S4(NO2)          9357 non-null   float64
10  PT08.S5(O3)           9357 non-null   float64
11  T                     9357 non-null   float64
12  RH                    9357 non-null   float64
13  AH                    9357 non-null   float64
dtypes: datetime64[ns](1), float64(12), int64(1)
memory usage: 1023.5 KB
```

```
In [5]: data.rename(columns = {'PT08.S4(NO2)': 'PT08_S4_NO2'}, inplace = True)
```

```
In [6]: data.corr()['PT08_S4_NO2']
```

```
Out[6]: CO(GT)                -0.073721
PT08.S1(CO)                0.845133
NMHC(GT)                   0.162689
C6H6(GT)                   0.774649
PT08.S2(NMHC)              0.874761
NOx(GT)                    0.035580
PT08.S3(NOx)               0.122672
NO2(GT)                   -0.022092
PT08_S4_NO2                1.000000
PT08.S5(O3)                0.723670
T                           0.755053
RH                          0.640685
AH                          0.691889
Name: PT08_S4_NO2, dtype: float64
```

The matrix of regression variables X will contain two variables:

- Temperature T
- Absolute Humidity AH

STEP 1: Prepare the data

```
In [7]: type(data['Date_Time'])
```

```
Out[7]: pandas.core.series.Series
```

```
In [8]: data['DateTimeIndex'] = pd.to_datetime(data['Date_Time'])  
data = data.set_index(keys=['DateTimeIndex'])
```

Set the frequency attribute of the index to Hourly. This will create several empty rows corresponding to the missing hourly measurements in the original data set. Fill up all the empty data cells with the mean of the corresponding column.

```
In [9]: ata = data.asfreq('H')  
data = data.fillna(data.mean(numeric_only=True))  
data.shape
```

```
Out[9]: (9357, 14)
```

```
In [10]: data.isin([np.nan, np.inf, -np.inf]).sum()
```

```
Out[10]: Date_Time      0  
CO(GT)      0  
PT08.S1(CO)  0  
NMHC(GT)    0  
C6H6(GT)    0  
PT08.S2(NMHC) 0  
NOx(GT)     0  
PT08.S3(NOx) 0  
NO2(GT)     0  
PT08.S4_NO2  0  
PT08.S5(O3)  0  
T           0  
RH          0  
AH          0  
dtype: int64
```

```
In [11]: dataset_len = len(data)
split_index = round(dataset_len*0.9)
train_set_end_date = data.index[split_index]
```

```
In [12]: df_train = data.loc[data.index <= train_set_end_date].copy()
df_test = data.loc[data.index > train_set_end_date].copy()
```

STEP 2: Create a Linear Regression model

```
In [13]: from patsy import dmatrices
expr = 'PT08_S4_NO2 ~ T + AH'
```

```
In [14]: y_train, X_train = dmatrices(expr, df_train, return_type='dataframe')
y_test, X_test = dmatrices(expr, df_test, return_type='dataframe')
```

Ordinary Least Squares Linear Regression model

```
In [15]: from statsmodels.regression import linear_model

olsr = linear_model.OLS(y_train, X_train).fit()
```

```
In [16]: olsr.summary()
```

Out[16]:

OLS Regression Results

Dep. Variable:	PT08_S4_NO2	R-squared:	0.654
Model:	OLS	Adj. R-squared:	0.654
Method:	Least Squares	F-statistic:	7969.
Date:	Tue, 31 May 2022	Prob (F-statistic):	0.00
Time:	17:23:49	Log-Likelihood:	-59426.
No. Observations:	8422	AIC:	1.189e+05
Df Residuals:	8419	BIC:	1.189e+05
Df Model:	2		
Covariance Type:	nonrobust		
	coef	std err	t P> t [0.025 0.975]
Intercept	1113.6844	7.218	154.302 0.000 1099.536 1127.833
T	20.5513	0.365	56.237 0.000 19.835 21.268
AH	-13.9624	0.405	-34.444 0.000 -14.757 -13.168
Omnibus:	904.777	Durbin-Watson:	0.286
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1410.196
Skew:	0.785	Prob(JB):	6.02e-307
Kurtosis:	4.247	Cond. No.	144.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression coefficients of both regression variables T and AH are significant at a 99.99% confidence level as indicated by their P values ($P > |t|$ column) which are essentially 0.


The second thing to note in these results is the output of the Durbin-Watson test which measures the degree of LAG-1 auto-correlation in the residual errors of regression. A value of 2 implies no LAG-1 auto-correlation. A value closer to 0 implies strong positive auto-correlation while a value close to 4 implies a strong negative auto-correlation at LAG-1 among the residuals errors ϵ .

In the above output, we see that the DW test statistic is 0.28 indicating a strong positive auto-correlation among the residual errors of regression at LAG-1. This was completely expected since the underlying data is a time series and the linear regression model has failed to explain the auto-correlation in the dependent variable. The DW test statistic just confirms it.

```
In [17]: import statsmodels.graphics.tsaplots as tsa

plt.rcParams.update({'figure.figsize':(10,5), 'figure.dpi':120})

tsa.plot_acf(olsr.resid, alpha=0.05)
plt.show()
```



The ACF tells us three things:

- There are strong auto-correlations extending out to multiple lags indicating that the residual errors time series has a trend. We'll need to de-trend this time series by using one or possibly 2 orders of differencing. Thus, the parameter d is likely to be 1, or possibly 2.
- The wavelike pattern in the ACF evidences a seasonal variation in the data.
- The peak at $\text{LAG} = 24$ indicates that the seasonal period is likely to be 24 hours. i.e. m is likely to be 24. This seems reasonable for data containing vehicular pollution measurements. We'll soon verify this guess using the time series decomposition plot.

Before we estimate the rest of the (S)ARIMA parameters, let's difference the time series once i.e. $d=1$:


```
In [18]: olsr.resid.head(2)
```

```
Out[18]: DateTimeIndex
2004-03-10 18:00:00    309.398174
2004-03-10 19:00:00    181.863043
dtype: float64
```

```
In [19]: olsr_resid_diff_1 = olsr.resid.diff()
olsr_resid_diff_1.head(2)
```

```
Out[19]: DateTimeIndex
2004-03-10 18:00:00         NaN
2004-03-10 19:00:00   -127.535131
dtype: float64
```


```
In [20]: olsr_resid_diff_1.isnull().sum()
```

```
Out[20]: 1
```

```
In [21]: olsr_resid_diff_1 = olsr_resid_diff_1.dropna()
```

```
In [22]: plt.rcParams.update({'figure.figsize':(10,5), 'figure.dpi':120})

tsa.plot_acf(olsr_resid_diff_1, alpha=0.05)
plt.show()
```




We now see a very different picture in the ACF. The auto-correlations are significantly reduced at all lags. The wavelike pattern still exists but that's because we did nothing to remove the possible seasonal variation. The LAG-24 auto-correlation is once again especially prominent.

We see that there is still a significant auto-correlation at LAG-1 in the differenced time series. We could try extinguishing it by taking one more difference, i.e. $d=2$ and plotting the resulting time series' ACF:

```
In [23]: olsr_resid_diff_2 = olsr_resid_diff_1.diff()
olsr_resid_diff_2 = olsr_resid_diff_2.dropna()

plt.rcParams.update({'figure.figsize':(10,5), 'figure.dpi':120})
```

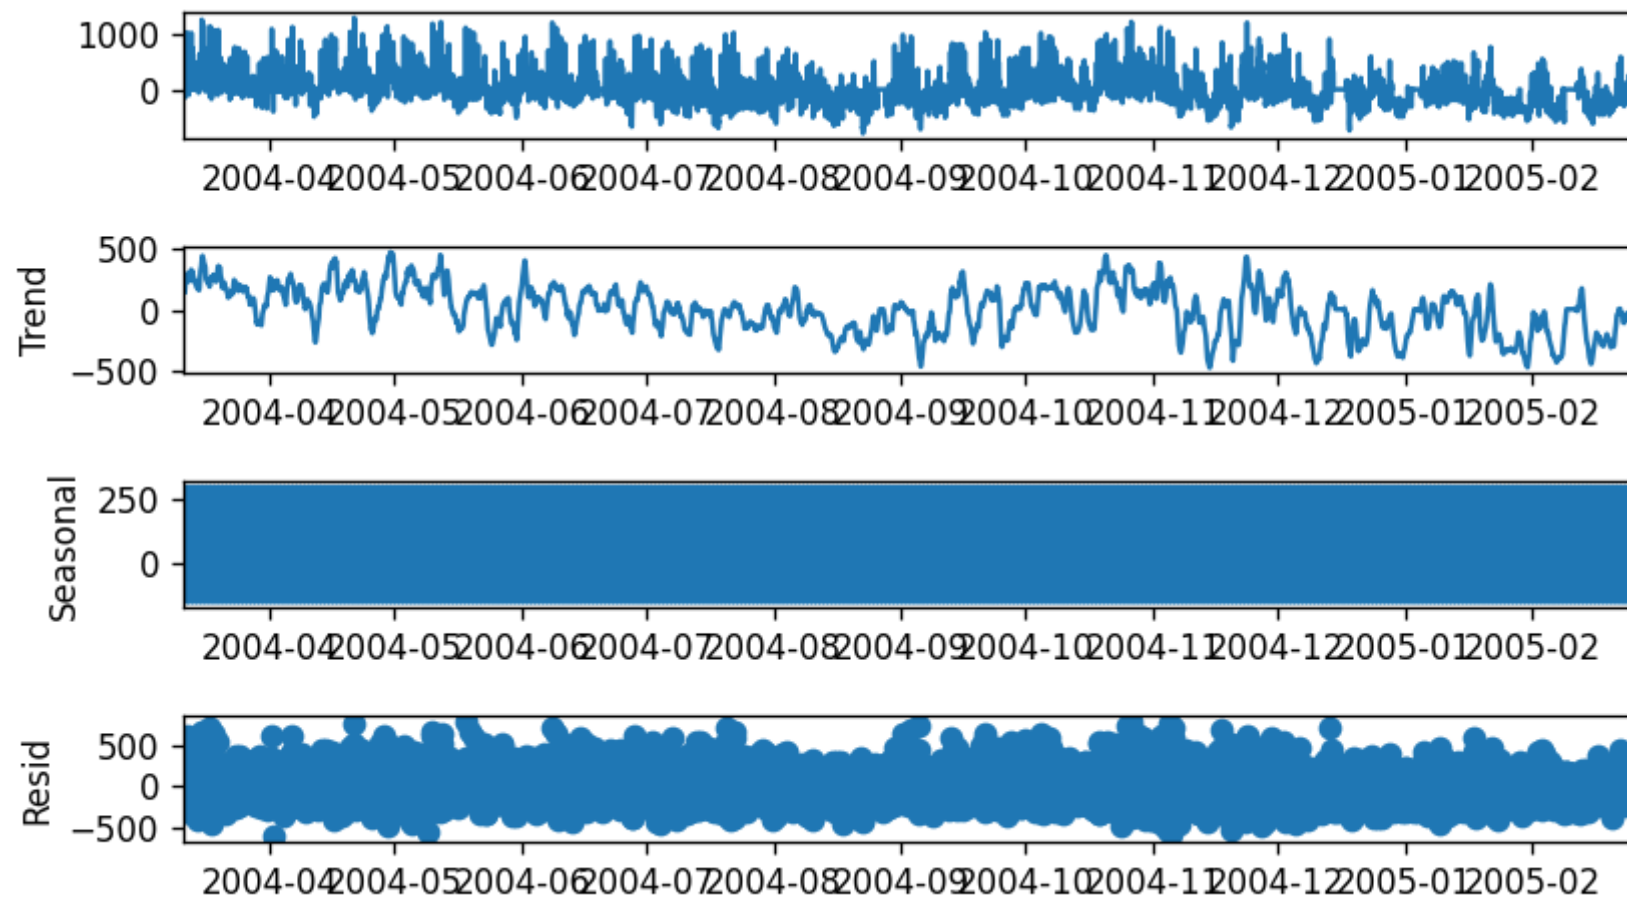
```
tsa.plot_acf(olsr_resid_diff_2, alpha=0.05)  
plt.show()
```




Differing the time series a second time has produced a heavy negative auto-correlation at LAG-1. This is bad sign. We seem to have over-done the differencing. We should stick with $d=1$.

```
In [24]: from statsmodels.tsa.seasonal import seasonal_decompose  
         components = seasonal_decompose(olsr.resid)  
  
         plt.rcParams.update({'figure.figsize':(7,4), 'figure.dpi':120})  
         components.plot()
```

Out[24]:






```
In [25]: olsr_resid_diff_1_24 = olsr_resid_diff_1.diff(periods=24)
olsr_resid_diff_1_24 = olsr_resid_diff_1_24.dropna()

plt.rcParams.update({'figure.figsize':(10,5), 'figure.dpi':100})

tsa.plot_acf(olsr_resid_diff_1_24, alpha=0.05)
plt.show()
```



The strong negative correlation at LAG-24 indicates a Seasonal MA (SMA) signature with order 1. i.e. $Q=1$. Moreover, an absence of positive correlation at LAG-1, indicates an absence of a Seasonal AR component. i.e. $P=0$.

We have fixed $P=0$, $D=1$ and $Q=1$, and $m=24$ hours

We have managed to estimate all 7 params of the SARIMA model as follows: $p=1$, $d=1$, $q=0$, $P=0$, $D=1$, $Q=1$ and $m=24$ i.e. $\text{SARIMAX}(1,1,0)$
 $(0,1,1)_{24}$

STEP 4: Build and fit the Regression Model with Seasonal ARIMA errors

```
In [26]: X_train_minus_intercept = X_train.drop('Intercept', axis=1)
X_train_minus_intercept = X_train_minus_intercept.asfreq('H')
```

```
In [27]: y_train = y_train.asfreq('H')
```

```
In [28]: from statsmodels.tsa.arima.model import ARIMA
sarimax_model = ARIMA(endog=y_train,
                      exog=X_train_minus_intercept,
                      order=(1,1,0),
                      seasonal_order=(0,1,1,24))
```

```
In [29]: sarimax_results = sarimax_model.fit()
```

```
In [30]: sarimax_results.summary()
```

Out[30]:

SARIMAX Results

Dep. Variable:	PT08_S4_NO2		No. Observations:	8422		
Model:	ARIMA(1, 1, 0)x(0, 1, [1], 24)			Log Likelihood	-51815.958	
Date:	Tue, 31 May 2022			AIC	103641.916	
Time:	17:24:53			BIC	103677.094	
Sample:	03-10-2004			HQIC	103653.928	
- 02-24-2005						
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
T	-4.6351	0.628	-7.383	0.000	-5.865	-3.405
AH	13.3304	0.673	19.808	0.000	12.011	14.649
ar.L1	-0.0041	0.008	-0.519	0.604	-0.020	0.011
ma.S.L24	-0.9184	0.003	-281.217	0.000	-0.925	-0.912
sigma2	1.334e+04	124.277	107.318	0.000	1.31e+04	1.36e+04
Ljung-Box (L1) (Q):	0.01	Jarque-Bera (JB):	5180.05			
Prob(Q):	0.93	Prob(JB):	0.00			
Heteroskedasticity (H):	0.66	Skew:	-0.10			
Prob(H) (two-sided):	0.00	Kurtosis:	6.84			

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [31]: sarimax_model = ARIMA(endog=y_train,
                                exog=X_train_minus_intercept,
                                order=(1,1,0),
                                seasonal_order=(0,1,0,24))
```



```
sarimax_results = sarimax_model.fit()
sarimax_results.summary()
```

Out[31]:

SARIMAX Results

Dep. Variable:		PT08_S4_NO2		No. Observations:		8422
Model:		ARIMA(1, 1, 0)x(0, 1, 0, 24)		Log Likelihood		-53887.658
Date:		Tue, 31 May 2022		AIC		107783.315
Time:		17:24:58		BIC		107811.458
Sample:		03-10-2004		HQIC		107792.925
		- 02-24-2005				
Covariance Type:		opg				
	coef	std err	z	P> z	[0.025	0.975]
T	-4.9003	0.688	-7.123	0.000	-6.249	-3.552
AH	13.6175	0.738	18.442	0.000	12.170	15.065
ar.L1	-0.0519	0.008	-6.482	0.000	-0.068	-0.036
sigma2	2.195e+04	208.195	105.422	0.000	2.15e+04	2.24e+04
Ljung-Box (L1) (Q):		0.67	Jarque-Bera (JB):		4288.52	
Prob(Q):		0.41	Prob(JB):		0.00	
Heteroskedasticity (H):		0.69	Skew:		0.06	
Prob(H) (two-sided):		0.00	Kurtosis:		6.50	

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

STEP 5: Prediction

```
In [32]: X_test_minus_intercept = X_test.drop('Intercept', axis=1)
X_test_minus_intercept = X_test_minus_intercept.asfreq('H')
y_test = y_test.asfreq('H')
```

```
In [33]: predictions = sarimax_results.get_forecast(steps=24, exog=X_test_minus_intercept[:24])
predictions.summary_frame().head()
```

```
Out[33]:
```

	PT08_S4_NO2	mean	mean_se	mean_ci_lower	mean_ci_upper
2005-02-24 16:00:00	1286.066400	148.149772	995.698183	1576.434617	
2005-02-24 17:00:00	1308.454485	204.153164	908.321635	1708.587335	
2005-02-24 18:00:00	1345.693115	248.033957	859.555492	1831.830738	
2005-02-24 19:00:00	1376.087728	285.232092	817.043100	1935.132356	
2005-02-24 20:00:00	1388.823150	318.110281	765.338456	2012.307845	

```
In [34]: plt.rcParams.update({'figure.figsize':(10,5), 'figure.dpi':100})


predicted, = plt.plot(X_test_minus_intercept[:24].index, predictions.summary_frame()['mean'], 'go-', label='Predicted')

actual, = plt.plot(y_test[:24], 'ro-', label='Actual')

lower, = plt.plot(X_test_minus_intercept[:24].index, predictions.summary_frame()['mean_ci_lower'], color='#990099', marker='.', label='Lower CI')

upper, = plt.plot(X_test_minus_intercept[:24].index, predictions.summary_frame()['mean_ci_upper'], color='#0000cc', marker='.', label='Upper CI')

plt.fill_between(X_test_minus_intercept[:24].index, predictions.summary_frame()['mean_ci_lower'], predictions.summary_frame()['mean_ci_upper'], color='lightblue')
plt.legend(handles=[predicted, actual, lower, upper])
plt.show()
```



ARIMA model

```
In [35]: df = pd.read_csv('https://raw.githubusercontent.com/selva86/datasets/master/wwwusage.csv', names=['value'], header=0)
```

Find the order of differencing (d) in ARIMA model

```
In [36]: from statsmodels.tsa.stattools import adfuller
from numpy import log
result = adfuller(df)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
```

ADF Statistic: -2.464240

p-value: 0.124419


```
In [37]: import numpy as np, pandas as pd
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
import matplotlib.pyplot as plt
plt.rcParams.update({'figure.figsize':(9,7), 'figure.dpi':120})

# Original Series
fig, axes = plt.subplots(3, 2, sharex=True)
axes[0, 0].plot(df.value); axes[0, 0].set_title('Original Series')
plot_acf(data[['PT08_S4_N02']], ax=axes[0, 1])

# 1st Differencing
axes[1, 0].plot(df.value.diff()); axes[1, 0].set_title('1st Order Differencing')
plot_acf(df.value.diff().dropna(), ax=axes[1, 1])

# 2nd Differencing
axes[2, 0].plot(df.value.diff().diff()); axes[2, 0].set_title('2nd Order Differencing')
plot_acf(df.value.diff().diff().dropna(), ax=axes[2, 1])

plt.show()
```



Auto Arima Forecast

```
In [38]: from statsmodels.tsa.arima_model import ARIMA
import pmdarima as pm
model = pm.auto_arima(df.value, start_p=1, start_q=1,
                      test='adf',          # use adftest to find optimal 'd'
                      max_p=3, max_q=3,    # maximum p and q
                      m=1,                 # frequency of series
                      d=None,              # let model determine 'd'
                      seasonal=False,      # No Seasonality
                      start_P=0,
                      D=0,
                      trace=True,
                      error_action='ignore',
                      suppress_warnings=True,
                      stepwise=True)

print(model.summary())
```

Performing stepwise search to minimize aic

```

ARIMA(1,2,1)(0,0,0)[0] intercept : AIC=525.587, Time=0.03 sec
ARIMA(0,2,0)(0,0,0)[0] intercept : AIC=533.474, Time=0.01 sec
ARIMA(1,2,0)(0,0,0)[0] intercept : AIC=532.437, Time=0.02 sec
ARIMA(0,2,1)(0,0,0)[0] intercept : AIC=525.893, Time=0.03 sec
ARIMA(0,2,0)(0,0,0)[0]          : AIC=531.477, Time=0.01 sec
ARIMA(2,2,1)(0,0,0)[0] intercept : AIC=515.248, Time=0.05 sec
ARIMA(2,2,0)(0,0,0)[0] intercept : AIC=513.459, Time=0.03 sec
ARIMA(3,2,0)(0,0,0)[0] intercept : AIC=515.284, Time=0.05 sec
ARIMA(3,2,1)(0,0,0)[0] intercept : AIC=inf, Time=0.22 sec
ARIMA(2,2,0)(0,0,0)[0]          : AIC=511.465, Time=0.03 sec
ARIMA(1,2,0)(0,0,0)[0]          : AIC=530.444, Time=0.00 sec
ARIMA(3,2,0)(0,0,0)[0]          : AIC=513.291, Time=0.03 sec
ARIMA(2,2,1)(0,0,0)[0]          : AIC=513.256, Time=0.03 sec
ARIMA(1,2,1)(0,0,0)[0]          : AIC=523.592, Time=0.02 sec
ARIMA(3,2,1)(0,0,0)[0]          : AIC=inf, Time=0.16 sec

```

Best model: ARIMA(2,2,0)(0,0,0)[0]

Total fit time: 0.714 seconds

SARIMAX Results

```

=====
Dep. Variable:          y      No. Observations:          100
Model:                SARIMAX(2, 2, 0)  Log Likelihood      -252.732
Date:                Tue, 31 May 2022    AIC                  511.465
Time:                17:25:01           BIC                  519.220
Sample:              0                HQIC                  514.601
                             - 100

```

Covariance Type: opg

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.2579      0.103      2.510      0.012      0.056      0.459
ar.L2         -0.4407      0.087     -5.093      0.000     -0.610     -0.271
sigma2        10.1268      1.519      6.668      0.000      7.150     13.103
=====

```

```

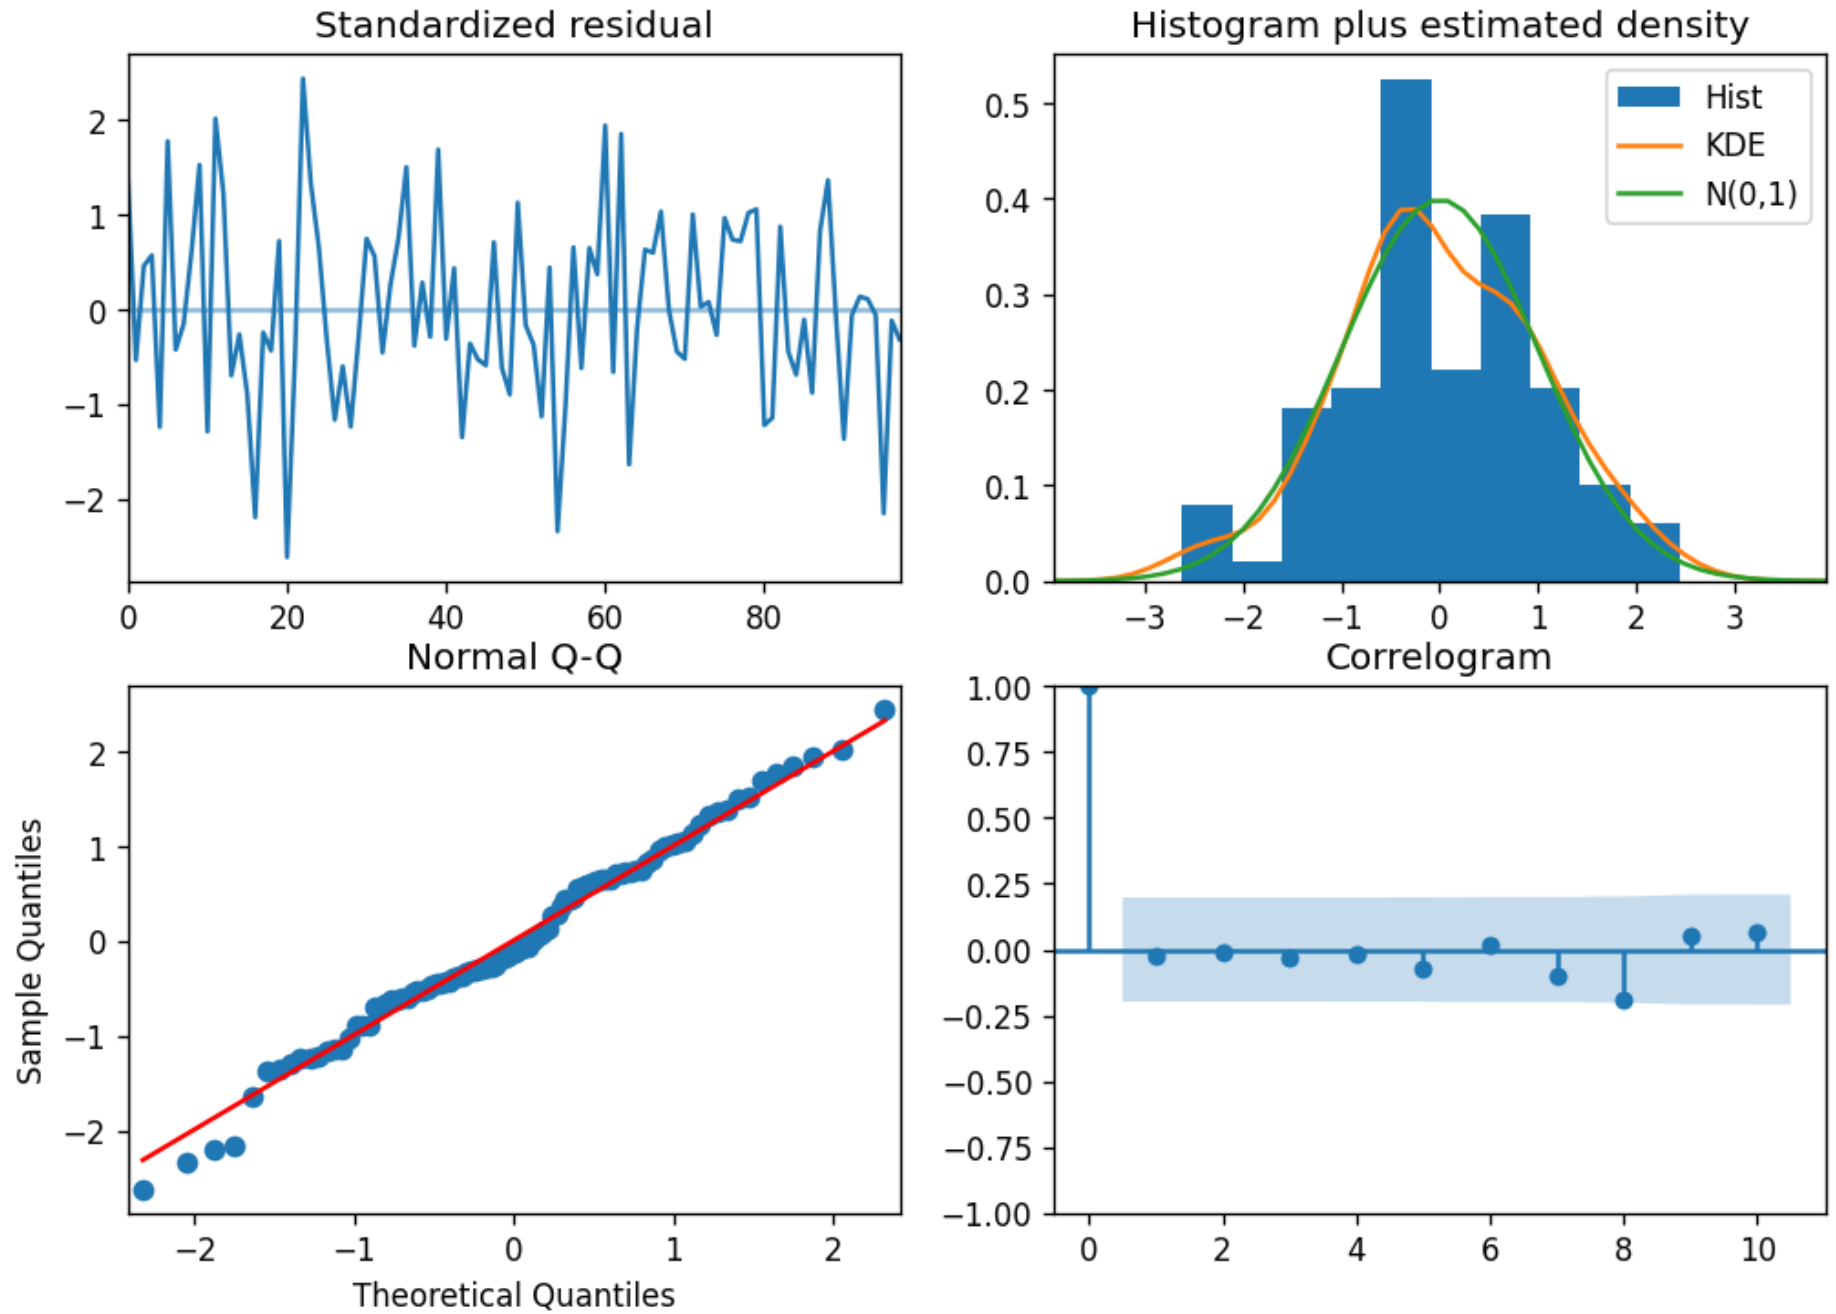
=====
Ljung-Box (L1) (Q):          0.05  Jarque-Bera (JB):          0.10
Prob(Q):                    0.82  Prob(JB):              0.95
Heteroskedasticity (H):      0.49  Skew:                 -0.07
Prob(H) (two-sided):        0.05  Kurtosis:             2.92
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [39]: model.plot_diagnostics(figsize=(10,7))  
plt.show()
```





```
In [40]: # Forecast
```

```
n_periods = 24
fc, confint = model.predict(n_periods=n_periods, return_conf_int=True)
index_of_fc = np.arange(len(df.value), len(df.value)+n_periods)

# make series for plotting purpose
fc_series = pd.Series(fc, index=index_of_fc)
lower_series = pd.Series(confint[:, 0], index=index_of_fc)
upper_series = pd.Series(confint[:, 1], index=index_of_fc)

# Plot
plt.plot(df.value)
plt.plot(fc_series, color='darkgreen')
plt.fill_between(lower_series.index,
                 lower_series,
                 upper_series,
                 color='k', alpha=.15)
plt.title("Final Forecast of WWW Usage")
plt.show()
```



In []: